

## Q14.6

March 16, 2024

**14.6** Consider the Protein data at the UCI machine learning repository called “Physiochemical Properties of Protein Tertiary Structure”. The data have 45,730 observations, one y-variable, RMSD, and nine x-variables. Use the first half of the data as training+validation data, and the remainder as test data. (We would usually randomize the order, but this setup allows a comparison on the same test data.) You can decide the details for validation. Make sure you standardize the x-variables.

- (a) Find a neural network to minimize squared error loss on the validation data. Explore a (small) number of hidden layers, the number of nodes within layers, the number of epochs, batch size, and learning rate. Summarize what helped and what did not help, and report on the best set of meta-parameters and the corresponding MSE for the validation data you found.
- (b) For the meta-parameters chosen in (a), report your MSE on the test data.
- (c) Add dropout after hidden layers (and adjust other meta-parameters as you see fit). Does dropout reduce the MSE on the validation set? If so, recompute the test error.

Loading libraries...

```
[ ]: import numpy as np
import sklearn as sk
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib as mpl
import torch
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
import math
import statistics
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

# change the device to mps to accelerate training
device = (
    "cuda"
```

```

    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")

```

```

[ ]: # read in the data
data = pd.read_csv("CASP.csv")
for i in data.columns[1:len(data.columns)]:
    # standardizing the variables
    data[i] = (data[i] - statistics.mean(data[i]))/statistics.stdev(data[i])

# extract the labels (first column)
labels = data["RMSD"]
# extract the data for prediction
df = data.loc[:, "F1": "F9"]
# convert the data to tensors
x_data = torch.tensor(data = df.values, dtype = torch.float32)
y_data = torch.tensor(data = labels.values, dtype = torch.float32)
# split into train/test
X_train = x_data[:len(data)//2]
X_test = x_data[len(data)//2:]
y_train = y_data[:len(data)//2]
y_test = y_data[len(data)//2:]
# split the training data into training and validation
# (assuming the ordering should still be maintained)
x_train, x_valid, y_train, y_valid = train_test_split(X_train, y_train,
    ↪test_size = 0.20, shuffle = False)

```

### 0.0.1 Part a)

Defining the neural network, “playing with layers” and meta-parameters.

```

[ ]: # set the number of epochs, batch size, and learning rate
epochs = 1000
batch_size = 8
lr = 1e-5

```

Defining the NN class.

Here, I have it set up to where you can input variable sizes, as will be demonstrated in the training loop.

```

[ ]: class SimpleNN(nn.Module):
    def __init__(self, input_size, hidden_sizes, output_size):
        super(SimpleNN, self).__init__()
        layers = [nn.Linear(input_size, hidden_sizes[0]), nn.ReLU()]

```

```

        for i in range(1, len(hidden_sizes)):
            layers.append(nn.Linear(hidden_sizes[i - 1], hidden_sizes[i]))
            layers.append(nn.ReLU())
        layers.append(nn.Linear(hidden_sizes[-1], output_size))
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)

```

```

[ ]: input_size = 9
     output_size = 1
     criterion = nn.MSELoss()
     # to test the input of hidden layers
     hidden_layers_config = [128]
     model = SimpleNN(input_size, hidden_layers_config, output_size).to(device)

```

```

[ ]: SimpleNN(
      (model): Sequential(
        (0): Linear(in_features=9, out_features=128, bias=True)
        (1): ReLU()
        (2): Linear(in_features=128, out_features=1, bias=True)
      )
)

```

```

[ ]: hidden_layers_configs = [[64], [128], [64, 64], [128, 128]]
     batch_sizes = [8, 16]
     learning_rates = [1e-4, 1e-5]
     epochs = 1000
     best_val_loss = float('inf')
     patience = 10
     best_model_params = None
     hyperparameters_log = []
     for hidden_layers in hidden_layers_configs:
         for batch_size in batch_sizes:
             for lr in learning_rates:
                 # reset the model
                 model = SimpleNN(input_size, hidden_layers, output_size).to(device)
                 # set the criterion and the optimizer
                 criterion = nn.MSELoss()
                 optimizer = optim.Adam(model.parameters(), lr=lr)

                 # create the data loader
                 train_loader = DataLoader(TensorDataset(x_train, y_train),
                 ↪ batch_size=batch_size, shuffle=True)
                 valid_loader = DataLoader(TensorDataset(x_valid, y_valid),
                 ↪ batch_size=batch_size, shuffle=False)

```

```

        local_best_val_loss = float('inf') # tracking the best lost for
↳the config
        epochs_no_improve = 0

    for epoch in range(epochs):
        model.train()
        for inputs, targets in train_loader:
            inputs, targets = inputs.to(device), targets.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs.squeeze(), targets)
            loss.backward()
            optimizer.step()

        # validation loop
        model.eval()
        with torch.no_grad():
            val_loss = 0
            for inputs, targets in valid_loader:
                inputs, targets = inputs.to(device), targets.to(device)
                outputs = model(inputs)
                val_loss += criterion(outputs.squeeze(), targets).item()
            val_loss /= len(valid_loader)

        if val_loss < local_best_val_loss:
            local_best_val_loss = val_loss
            if val_loss < best_val_loss:
                best_val_loss = val_loss
                best_model_params = {
                    "hidden_layers": hidden_layers,
                    "batch_size": batch_size,
                    "learning_rate": lr,
                }
        else:
            epochs_no_improve += 1 # increment
            if epochs_no_improve >= patience:
                print(f"Early stopping triggered after {epoch + 1}
↳epochs.")
                break # early stop

        print(f"Epoch {epoch+1}/{epochs}, Training Loss: {loss.item()},
↳Validation Loss: {val_loss}")

    # record parameters
    hyperparameters_log.append({
        "params": {
            "hidden_layers": hidden_layers,

```

```

        "batch_size": batch_size,
        "learning_rate": lr,
    },
    "val_loss": local_best_val_loss
})

# printing the best parameters
print(f"Best model parameters: {best_model_params}, with validation loss:␣
↪{best_val_loss}")

```

So, our best model has 2 hidden layers, with 128 nodes, a batch size of 8, and a learning rate of 0.0001.

This leads to an  $MSE_{\text{validation}} = 19.914560827237743$ .

Since we don't save the model state, we retrain the model with the best parameters, this time going through every single epoch.

```

[ ]: # create the data loader
batch_size = best_model_params["batch_size"]
hidden_layers = best_model_params["hidden_layers"]
lr = best_model_params["learning_rate"]
model = SimpleNN(input_size, hidden_layers, output_size).to(device)
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=lr)

train_loader = DataLoader(TensorDataset(x_train, y_train),␣
↪batch_size=batch_size, shuffle=True)
valid_loader = DataLoader(TensorDataset(x_valid, y_valid),␣
↪batch_size=batch_size, shuffle=False)

for epoch in range(epochs):
    model.train()
    for inputs, targets in train_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs.squeeze(), targets)
        loss.backward()
        optimizer.step()

    # validation loop
    model.eval()
    with torch.no_grad():
        val_loss = 0
        for inputs, targets in valid_loader:
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = model(inputs)

```

```

        val_loss += criterion(outputs.squeeze(), targets).item()
    val_loss /= len(valid_loader)
    print(f"Epoch {epoch+1}/{epochs}, Training Loss: {loss.item()}, Validation_
↪Loss: {val_loss}")

```

So, after “fully training” the model using the best parameters as detailed above, we get  $MSE_{\text{validation}} = 15.49317671191859$ .

### 0.0.2 Part b)

Evaluating on the test data.

```

[ ]: # create the test data set (since the training and validation are defined in
↪the training loop)
# due to the variable batch size depending on performance
test_dataset = TensorDataset(X_test, y_test)
test_loader = DataLoader(dataset = test_dataset, batch_size = batch_size,
↪shuffle = False)

```

```

[ ]: # initialize a list to store predictions and true labels
predictions = []
true_labels = []

with torch.no_grad():
    test_loss = 0
    for inputs, targets in test_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        # forward pass
        outputs = model(inputs)
        # calculate the loss
        loss = criterion(outputs.squeeze(), targets)
        test_loss += loss.item()
        # store preds/true values
        outputs = outputs.squeeze()
        if outputs.ndim == 0:
            predictions.append(outputs.item())
            true_labels.append(targets.item())
        else:
            predictions.extend(outputs.tolist())
            true_labels.extend(targets.tolist())

# calculate mse on the test data
test_loss /= len(test_loader)

print(f"Mean Squared Error on Test Data: {test_loss}")

```

As we can see, we obtain  $MSE_{\text{test}} = 16.434121215966833$ .

### 0.0.3 Part c)

Adding a dropout layer.

The choice of  $p = 0.25$  for the dropout layer is arbitrary here.

```
[ ]: class SimpleNNDropout(nn.Module):
    def __init__(self, input_size, hidden_sizes, output_size, dropout_rate=0.
    ↪25):
        super(SimpleNNDropout, self).__init__()
        layers = [nn.Linear(input_size, hidden_sizes[0]), nn.ReLU(), nn.
    ↪Dropout(dropout_rate)]
        for i in range(1, len(hidden_sizes)):
            layers.append(nn.Linear(hidden_sizes[i - 1], hidden_sizes[i]))
            layers.append(nn.ReLU())
            layers.append(nn.Dropout(dropout_rate)) # add dropout after each
    ↪activation
        layers.append(nn.Linear(hidden_sizes[-1], output_size))
        self.model = nn.Sequential(*layers)

    def forward(self, x):
        return self.model(x)
```

```
[ ]: model = SimpleNNDropout(input_size, hidden_layers, output_size, dropout_rate =
    ↪0.25).to(device)
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=lr)
train_loader = DataLoader(TensorDataset(x_train, y_train),
    ↪batch_size=batch_size, shuffle=True)
valid_loader = DataLoader(TensorDataset(x_valid, y_valid),
    ↪batch_size=batch_size, shuffle=False)
for epoch in range(epochs):
    model.train()
    for inputs, targets in train_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs.squeeze(), targets)
        loss.backward()
        optimizer.step()

    # validation loop
    model.eval()
    with torch.no_grad():
        val_loss = 0
        for inputs, targets in valid_loader:
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = model(inputs)
```

```

        val_loss += criterion(outputs.squeeze(), targets).item()
    val_loss /= len(valid_loader)
    print(f"Epoch {epoch+1}/{epochs}, Training Loss: {loss.item()}, Validation_
↪Loss: {val_loss}")

```

As we can see, we do get a marginal improvement in MSE:  $MSE_{\text{dropout validation}} = 14.936854807646958$ .

So, we compute the training error using dropout:

```

[ ]: # create the test data set (since the training and validation are defined in
↪the training loop)
# due to the variable batch size depending on performance
test_dataset = TensorDataset(X_test, y_test)
test_loader = DataLoader(dataset = test_dataset, batch_size = batch_size,
↪shuffle = False)

# initialize a list to store predictions and true labels
predictions = []
true_labels = []

with torch.no_grad():
    test_loss = 0
    for inputs, targets in test_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        # forward pass
        outputs = model(inputs).squeeze()
        # calculate the loss
        loss = criterion(outputs.squeeze(), targets.squeeze())
        test_loss += loss.item()
        # store preds/true values
        outputs = outputs.squeeze()
        if outputs.ndim == 0:
            predictions.append(outputs.item())
            true_labels.append(targets.item())
        else:
            predictions.extend(outputs.tolist())
            true_labels.extend(targets.tolist())

# calculate mse on the test data
test_loss /= len(test_loader)

print(f"Mean Squared Error on Test Data: {test_loss}")

```

So, we get  $MSE_{\text{dropout test}} = 15.5339203014704$  - an improvement over the previous MSE of 16.434121215966833.